

A close-up photograph of a hand holding a piece of brown sandpaper and sanding a piece of light blue painted wood. The wood grain is visible, and the sandpaper is being moved across the surface. The background is a blurred wooden surface.

Lecture 6: Smoothing and Model Complexity in Language Models

Instructor: Jackie CK Cheung

COMP-550

J&M Ch 6.3 (1st ed); J&M Ch 4.5 (2nd ed) ;

J&M Ch 3.3 – 3.4 (3rd ed)

Outline

Justification of MLE probabilistically

Overfitting and unseen data

Dealing with unseen data

- Interpolation
- Good-Turing smoothing

Language Modelling

Predict the next word given some context

Mary had a little _____

- *lamb* GOOD
- *accident* GOOD?
- *very* BAD
- *up* BAD

N-grams

Make a **conditional independence assumption** to make the job of learning the probability distribution easier.

- Context = the previous N-1 words

Common choices: N is between 1 and 3

Unigram model

$$P(w_N|C) = P(w_N)$$

Bigram model

$$P(w_N|C) = P(w_N|w_{N-1})$$

Trigram model

$$P(w_N|C) = P(w_N|w_{N-1}, w_{N-2})$$

Deriving Parameters from Counts

Simplest method: count N-gram frequencies, then divide by the total count

e.g.,

Unigram: $P(\text{cats}) = \text{Count}(\text{cats}) / \text{Count}(\text{all words in corpus})$

Bigram: $P(\text{cats} \mid \text{the}) = \text{Count}(\text{the cats}) / \text{Count}(\text{the})$

Trigram: $P(\text{cats} \mid \text{feed the}) = ?$

These are the **maximum likelihood estimates (MLE)**.

Why MLE?

Why is using relative frequencies of events in the training set a reasonable thing to do?

Why is it called maximum likelihood estimation?

Are there other choices for estimating model parameters?

What is Maximum Likelihood?

Find the model parameters that maximize the likelihood of (i.e., the probability of generating) the training corpus, X .

Find θ^{MLE} s.t. $P(X; \theta^{MLE})$ is maximized

In language modelling:

- Words are random variables that are drawn from a categorical probability distribution
- For a given context, they are i.i.d. (independently, and identically distributed)

Categorical Random Variables

1-of-K discrete outcomes, each with some probability

e.g., coin flip, die roll, draw a word from a language model

Probability of a training corpus, $C = x_1, x_2, \dots, x_N$:

$$\begin{aligned} K = 2: P(C; \theta) &= \prod_{n=1}^N P(x_n; \theta) \\ &= \theta^{N_1} (1 - \theta)^{N_0} \end{aligned}$$

Can similarly extend for $K > 2$

Notes:

- When $K=2$, it is called a Bernoulli distribution
- Sometimes incorrectly called a multinomial distribution, which is something else

Maximizing Quantities

Calculus to the rescue!

Take derivative and set to 0.

Trick: maximize the log likelihood instead (math works out better)

- This is okay because log is a monotonically increasing function, so it doesn't affect the location of the maximum.

MLE Derivation for a Bernoulli

Maximize the log likelihood:

$$\begin{aligned}\log P(C; \theta) &= \log(\theta^{N_1} (1 - \theta)^{N_0}) \\ &= N_1 \log \theta + N_0 \log(1 - \theta)\end{aligned}$$

$$\begin{aligned}\frac{d}{d\theta} \log P(C; \theta) &= \frac{N_1}{\theta} - \frac{N_0}{1 - \theta} = 0 \\ \frac{N_1}{\theta} &= \frac{N_0}{1 - \theta}\end{aligned}$$

Solve this to get:

$$\theta = \frac{N_1}{N_0 + N_1}$$

Or,

$$\theta = \frac{N_1}{N}$$

MLE Derivation for a Categorical

The above generalizes to the case where $K > 2$.

Parameters are now $\theta_0, \theta_1, \theta_2, \dots, \theta_{K-1}$

Counts are now $N_0, N_1, N_2, \dots, N_{K-1}$

Note: Need to add a constraint that $\sum_{i=0}^{K-1} \theta_i = 1$ to ensure that the parameters specify a probability distribution.

Use the method of *Lagrange multipliers*

Do this derivation in your next study group!

Hint:
$$L(\theta) = \log P(C; \theta) - \lambda(\sum_{i=0}^{K-1} \theta_i - 1)$$

$$L(\theta) = \sum_i N_i \log \theta_i - \lambda(\sum_i \theta_i - 1)$$

Find partial derivatives $\partial L / \partial \theta_i$ and $\partial L / \partial \lambda$, set them to zero and solve resulting system of equations

Steps

1. Gather a large, representative training corpus
2. Learn the parameters from the corpus to build the model
3. **Once the model is fixed, use the model to evaluate on testing data**

Overfitting

MLE often gives us a model that is too good of a fit to the training data. This is called **overfitting**.

- Words that we haven't seen
- The probabilities of the words and N-grams that we have seen are not representative of the true distribution.

But when testing, we evaluate the LM *on unseen data*. Overfitting lowers performance.

Out Of Vocabulary (OOV) Items

Suppose we train a LM on the WSJ corpus, which is about economic news in 1987 – 1989. What probability would be assigned to *Brexit*? To *QAnon*?

In general, we know that there will be many words in the test data that are not in the training data, no matter how large the training corpus is.

- Neologisms, typos, parts of the text in foreign languages, etc.
- Remember Zipf's law and the long tail

Explicitly Modelling OOV Items

During training:

- Pick some frequency threshold
- All vocabulary items that occur less frequently are replaced by an <UNK>
- Now, treat <UNK> as another vocabulary item

During testing:

- Any unseen words are called <UNK>

Smoothing

Training corpus does not have all the words

- Add a special <UNK> symbol for unknown words

Estimates for infrequent words are unreliable

- Modify our probability distributions

Smoothe the probability distributions to shift probability mass to cases that we haven't seen before or are unsure about

Prior Beliefs About Parameters

Smoothing means we are no longer doing MLE. We now have some **prior belief** about what the parameters should be like:

MLE:

Find θ^{MLE} s.t. $P(X; \theta^{MLE})$ is maximized

With smoothing:

Find θ^{smooth} s.t. $P(X; \theta^{smooth})P(\theta^{smooth})$ is maximized

Add- δ Smoothing

Modify our estimates by adding a certain amount to the frequency of each word. (sometimes called **pseudocounts**)

e.g., unigram model

$$P(w) = \frac{\text{Count}(w) + \delta}{N + \delta * |\text{Lexicon}|}$$

where N is the corpus size in #tokens

Pros: simple

Cons: not the best approach; how to pick δ ? Depends on sizes of lexicon and corpus

When $\delta = 1$, this is called **Laplace discounting**

Add- δ Smoothing in General

For a categorical distribution with k possible outcomes and a training corpus of C trials:

$$P(outcome) = \frac{\text{Count}(outcome) + \delta}{C + \delta * k}$$

So, what would be the formula for a bigram model, $P(w_t|w_{t-1})$?

Example

Suppose we have a LM with a vocabulary of 20,000 items.

In the training corpus, we see donkey 10 times.

- Of these, in 5 times it was followed by the word kong.
- In the other 5 times, it was followed by another word.

What is the MLE estimate of $P(\textit{kong} | \textit{donkey})$?

What is the Laplace estimate of $P(\textit{kong} | \textit{donkey})$?

Example

Suppose we have a LM with a vocabulary of 20,000 items.

In the training corpus, we see donkey 10 times.

- Of these, in 5 times it was followed by the word kong.
- In the other 5 times, it was followed by another word.

What is the MLE estimate of $P(\textit{kong} | \textit{donkey})$?

5/10

What is the Laplace estimate of $P(\textit{kong} | \textit{donkey})$?

6/20010

Interpolation

In an N-gram model, as N increases, data sparsity (i.e., unseen or rarely seen events) becomes a bigger problem.

In an **interpolation**, use a lower N to mitigate the problem.

Simple Interpolation

e.g., combine trigram, bigram, unigram models

$$\begin{aligned}\hat{P}(w_t | w_{t-2}, w_{t-1}) = & \lambda_1 P^{MLE}(w_t | w_{t-2}, w_{t-1}) \\ & + \lambda_2 P^{MLE}(w_t | w_{t-1}) \\ & + \lambda_3 P^{MLE}(w_t)\end{aligned}$$

Need to set $\sum_i \lambda_i = 1$ so that the overall sum is a probability distribution

How to select λ_i ? Compare different values on your development set!

Good-Turing Smoothing

A more sophisticated method of modelling unseen events (usually N-grams)

Remember Zipf's lessons

- We shouldn't adjust all words/N-grams uniformly.
- The frequency of a word or N-gram is related to its rank—we should be able to model this!
- Unseen N-grams should behave a lot like N-grams that only occur once in a corpus
- N-grams that occur a lot should behave like other N-grams that occur a lot.

Count of Counts

Let's build a histogram to count how many events occur a certain number of times in the corpus.

Event frequency	# events with that frequency
1	$f_1 = 3993$
2	$f_2 = 1292$
3	$f_3 = 664$
...	...

- For some event in bin f_c , that event occurred c times in the corpus; c is the numerator in the MLE.
- **Idea:** re-estimate c using f_{c+1}

Good-Turing Smoothing Defined

Let N be total number of observed event-tokens, w_c be an event that occurs c times in the training corpus.

$$N = \sum_i f_i \times i$$

$$P(UNK) = f_1 / N \quad \leftarrow \text{Note that } P(UNK) \text{ is for all unseen events}$$

$$\text{Then: } c^* = \frac{(c+1)f_{c+1}}{f_c}$$

$$P(w_c) = c^* / N \quad \leftarrow \text{Note that } P(w_c) \text{ is for one event that occurs } c \text{ times}$$

Example:

Let N be 100,000.

Word frequency	# word-types
1	$f_1 = 3,993$
2	$f_2 = 1,292$
3	$f_3 = 664$
...	...

$$\begin{aligned} P(UNK) &= 3993 / 100000 \\ &= 0.03993 \\ &\text{(for all unseen events)} \\ &\text{(Divide by total estimated} \\ &\text{number of unseen word} \\ &\text{types.)} \end{aligned}$$

$$\begin{aligned} c_1^* &= 2 * 1292 / 3993 \\ &= 0.647 \end{aligned}$$

$$\begin{aligned} c_2^* &= 3 * 664 / 1292 \\ &= 1.542 \end{aligned}$$

Exercises

Suppose we have the following counts:

Word	ship	pass	camp	frock	soccer	mother	tops
Freq	5	3	2	2	1	1	1

Give the MLE and Good-Turing estimates for the probabilities of:

- any unknown word
- *soccer*
- *camp*
- *pass*

$$N = \sum_i f_i \times i$$

$$c^* = \frac{(c+1)f_{c+1}}{f_c}$$

$$P(UNK) = f_1 / N$$

$$P(w_c) = c^* / N$$

Good-Turing Refinement

As seen above, simple Good-Turing fails for higher values of c , because f_{c+1} is often 0.

Solution: Estimate f_c as a function of c

- We'll assume that a linear relationship exists between $\log c$ and $\log f_c$
- Use linear regression to learn this relationship:
$$\log f_c^{LR} = a \log c + b$$
- For lower values of c , we continue to use f_c ; for higher values of c , we use our new estimate f_c^{LR} .

Model Selection

We now have very many slightly different versions of the model (with different **hyperparameters**). How to decide between them?

Recall: use a **development / validation set**

Procedure:

1. Train one or more models on the training set
2. Test (repeatedly, if necessary) on the dev/val set; choose a final hyperparameter setting/model
3. Test the final model on the final testing set (once only)

Steps 1 and 2 can be structured using cross-validation

Model Complexity Trade-Offs

In general, there is a trade-off between:

- model expressivity; i.e., what trends you could capture about your data with your model
- how prone the model is to overfitting

If you use a highly expressive model (e.g., high values of N in N -gram modelling), it is much easier to overfit, and you need to do smoothing. OTOH, if your model is too weak, your performance will suffer as well.